One Step Ahead

Detecting Unusual Human Motions for QualityMinds

qualityminds

CONSTRUCTOR LEARNING

The Team

Jonas Voßemer

Ph.D. Sociology

Alaa Elshorbagy

Ph.D. Mathematics

Vincent v. Zitzewitz

Industrial Engineer

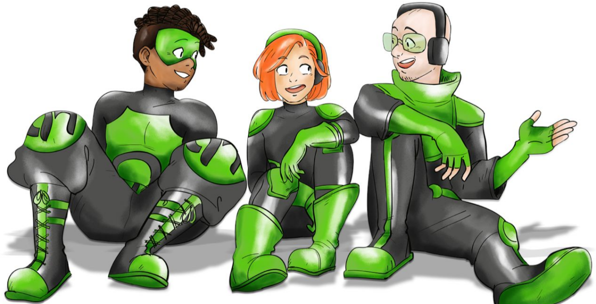

HMP & QM Business Case

With Whom …

… we collaborated

QualityMinds

→ Agile **IT Consultancy**

→ **Data Science** Branch

→ **Project:** Human motion prediction for Autonomous Driving

qualityminds

Why …

… Human Motion Prediction?

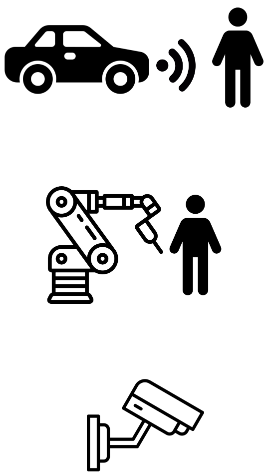

What …

… is Human Motion Prediction?

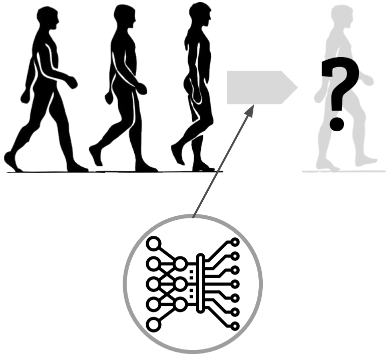

→ **Example** scenarios: eating, walking, phoning

How …

… we improve Human Motion Prediction?

Our Task:

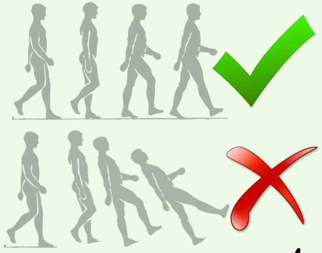

→ Quantify **"Anomaly"**

→ Unusual motions are **hard to predict**

Our Approach

Dataset → 3.6 million human poses

Data

180k sequences

35 frames ~ 1.4 seconds

32 joints

Position in 3D

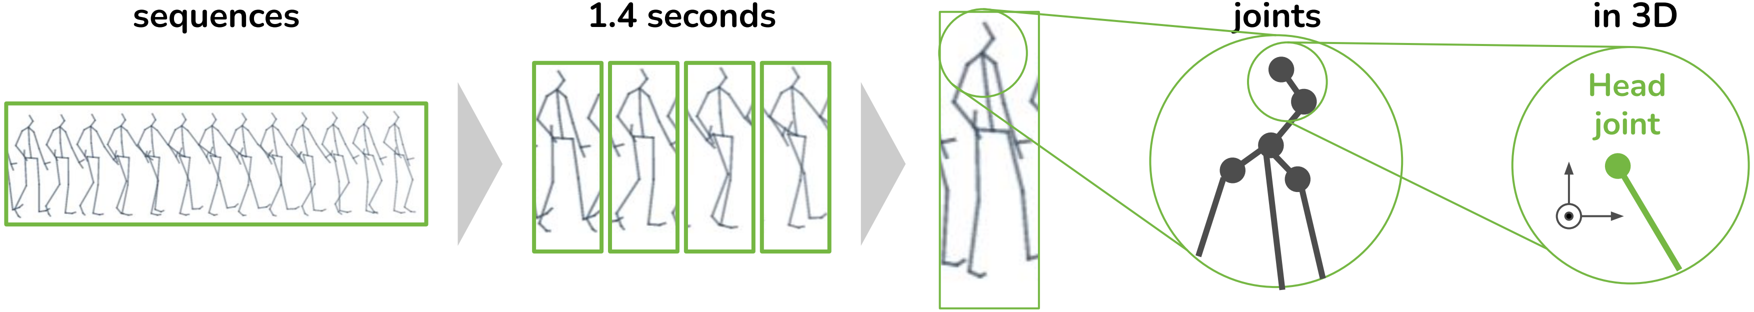

Our Approach

Methodology

Huge dataset → Dimen-sionality Reduction → **Simple dataset** → Scaling → **Scaled dataset** → Outliers Detection Models → **Quantification of “anomaly”**

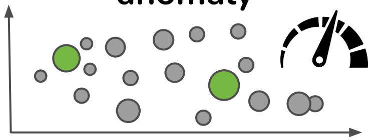

Validation

Prediction Error — **?** — **Decision Score**

Results: Outlier Detection App

Outlier Detection App

Select Your Analysis

Set Decision Score Threshold

0.70 0.75 0.80 0.85 0.90 0.95

Select Outlier Detection Model:

CBLOF ×

Color by:

Decision Score Threshold ×

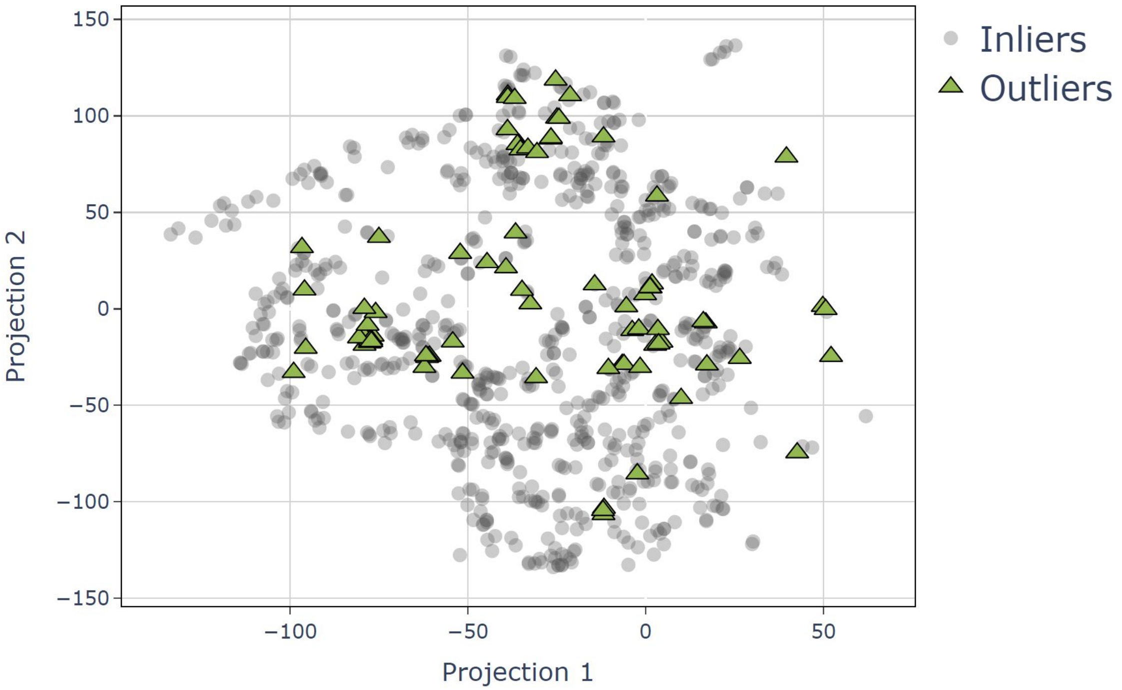

Results: Outlier Detection App

Outlier Detection App

Select Your Analysis

Set Decision Score Threshold

0.70 0.75 0.80 0.85 0.90 0.95

Select Outlier Detection Model:

CBLOF

Color by:

Decision Score Threshold

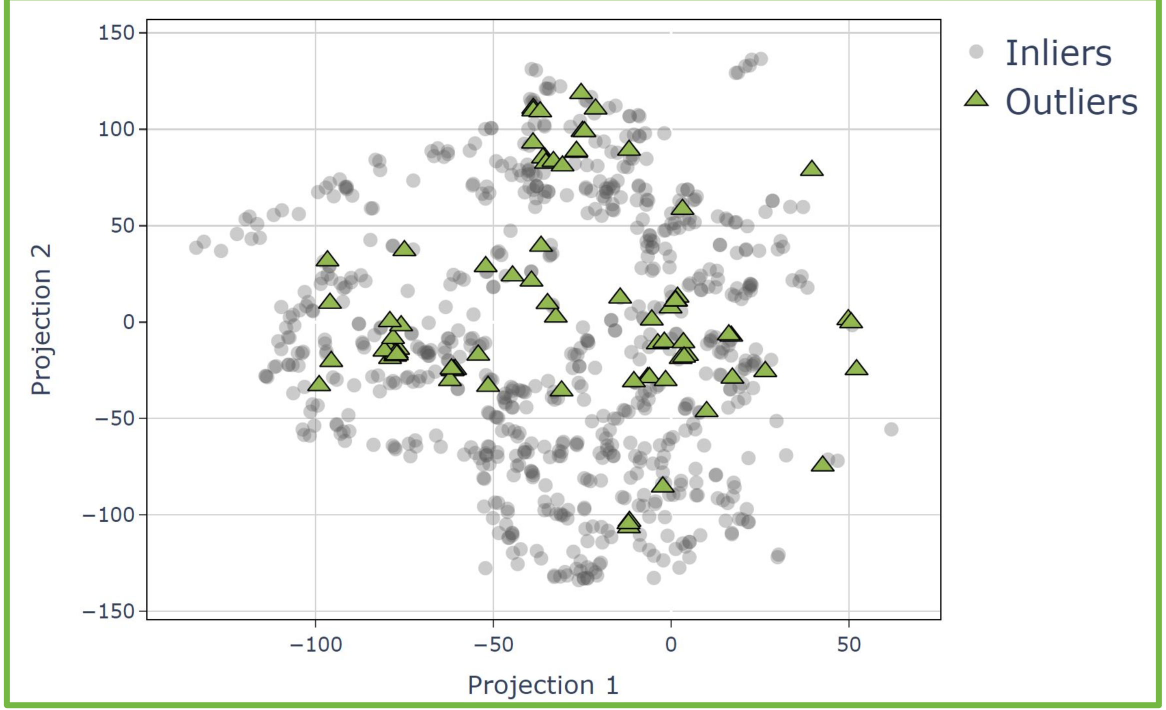

Results: Outlier Detection App

Outlier Detection App

Select Your Analysis

Set Decision Score Threshold

0.70 0.75 0.80 0.85 0.90 0.95

Select Outlier Detection Model:

CBLOF

Color by:

Decision Score Threshold

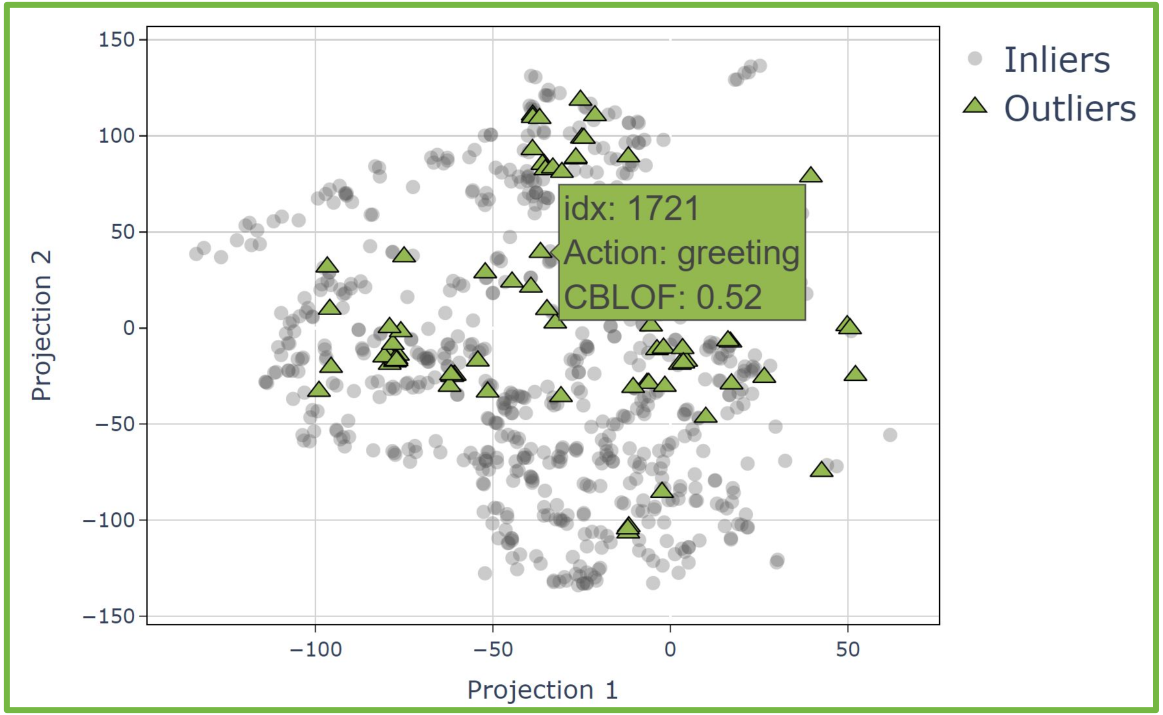

Results: Outlier Detection App

Outlier Detection App

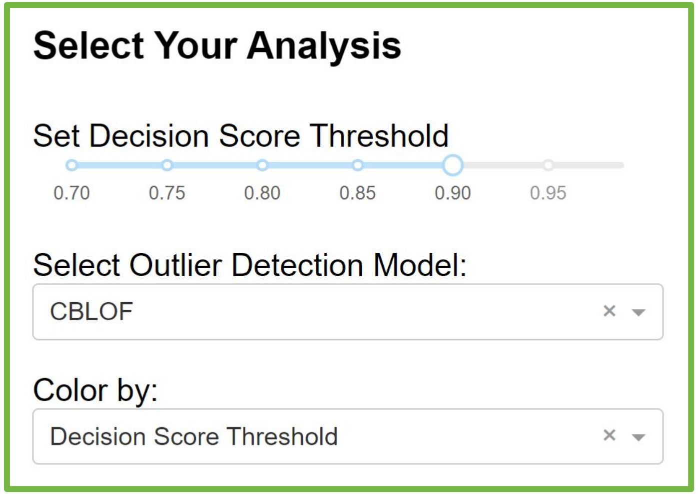

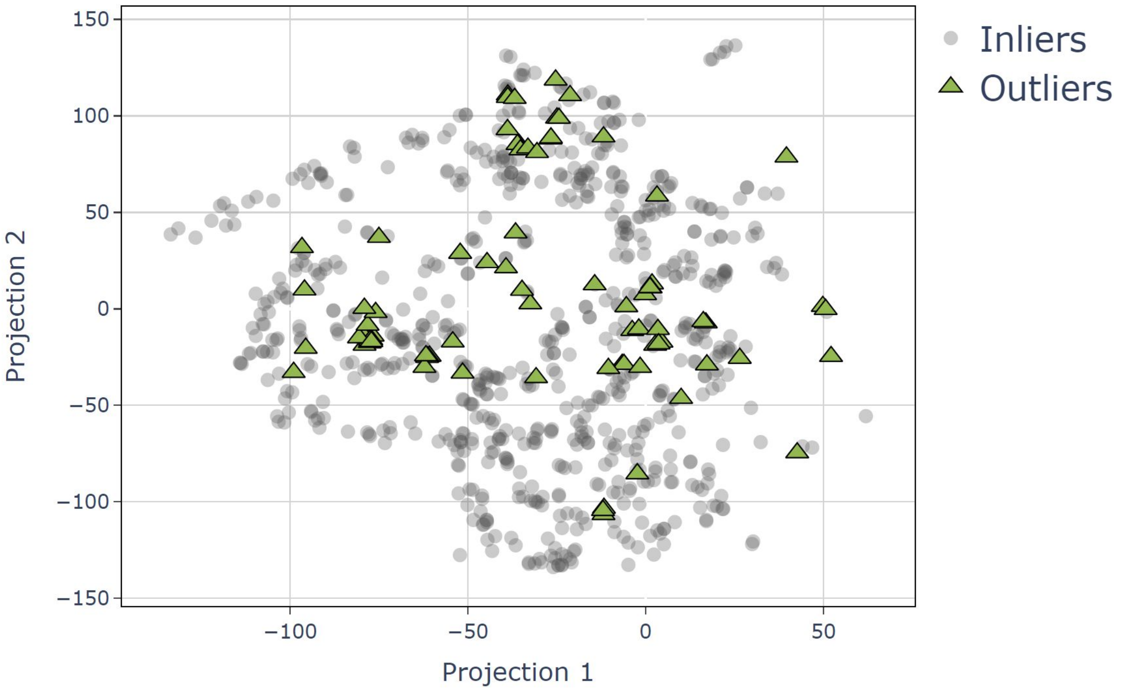

Results: Outlier Detection App

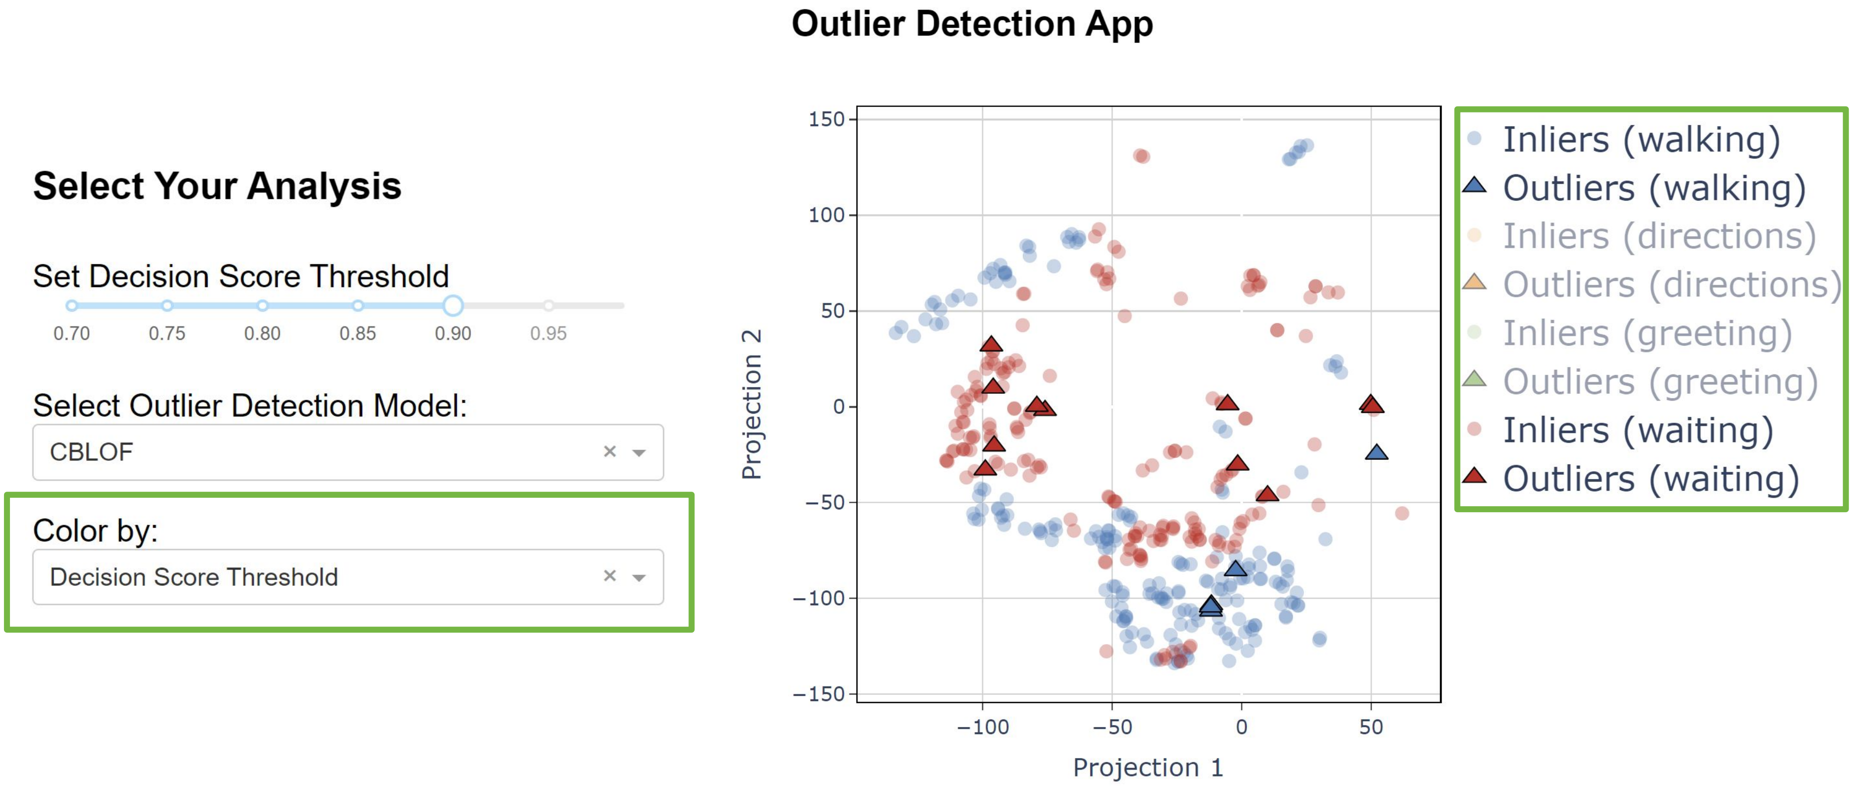

Results: Outlier Validation App

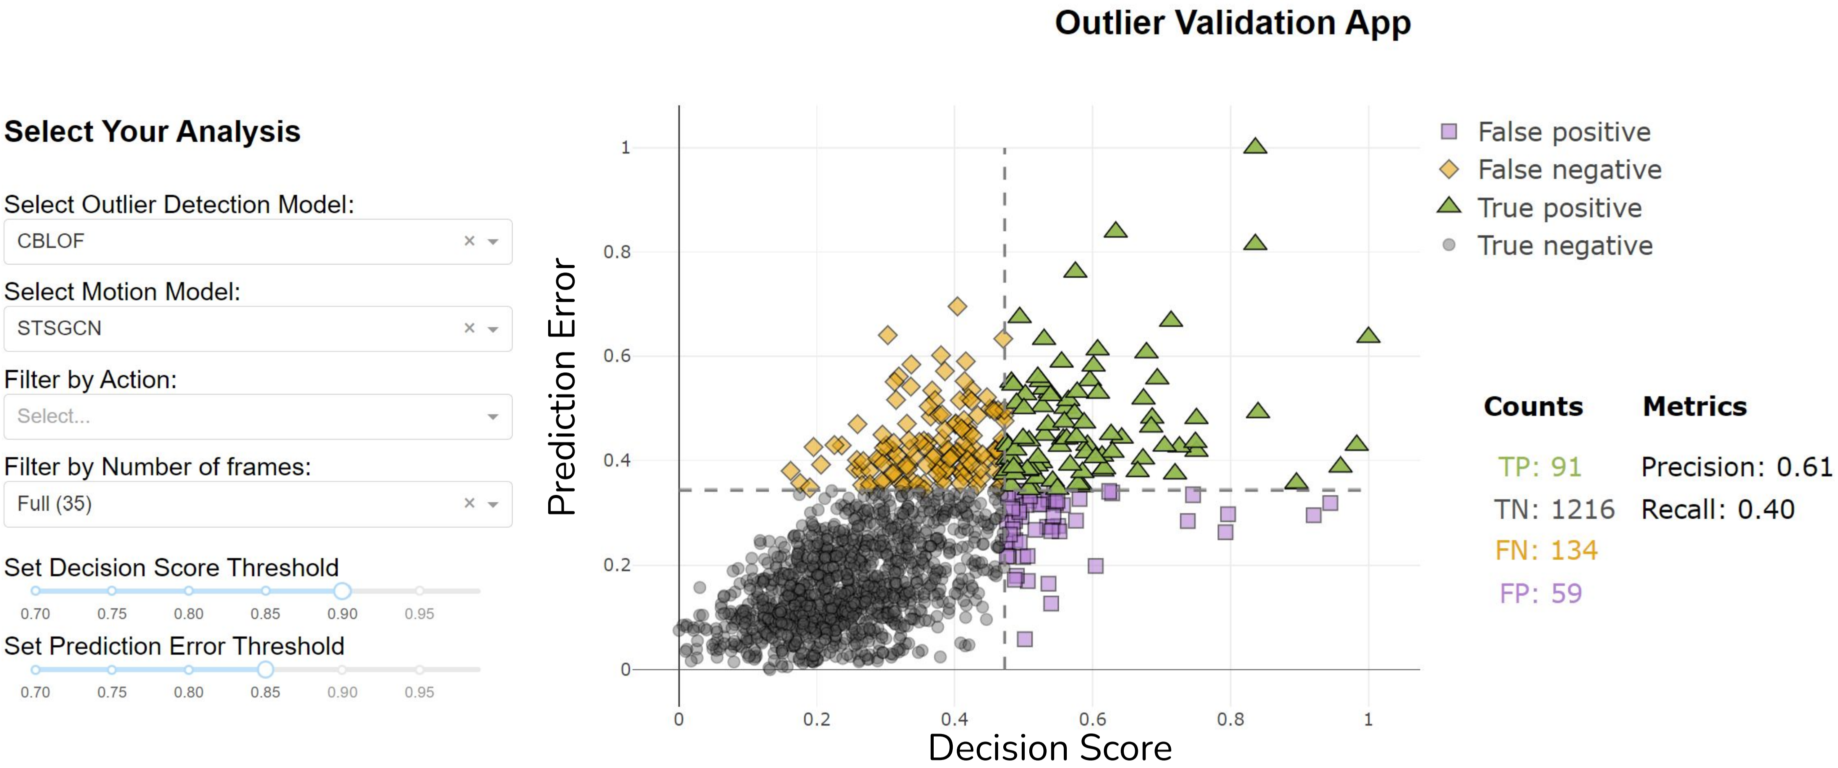

Results: Outlier Validation App

Outlier Validation App

Select Your Analysis

Select Outlier Detection Model:

CBLOF

Select Motion Model:

STSGCN

Filter by Action:

Select...

Filter by Number of frames:

Full (35)

Set Decision Score Threshold

0.70 0.75 0.80 0.85 0.90 0.95

Set Prediction Error Threshold

0.70 0.75 0.80 0.85 0.90 0.95

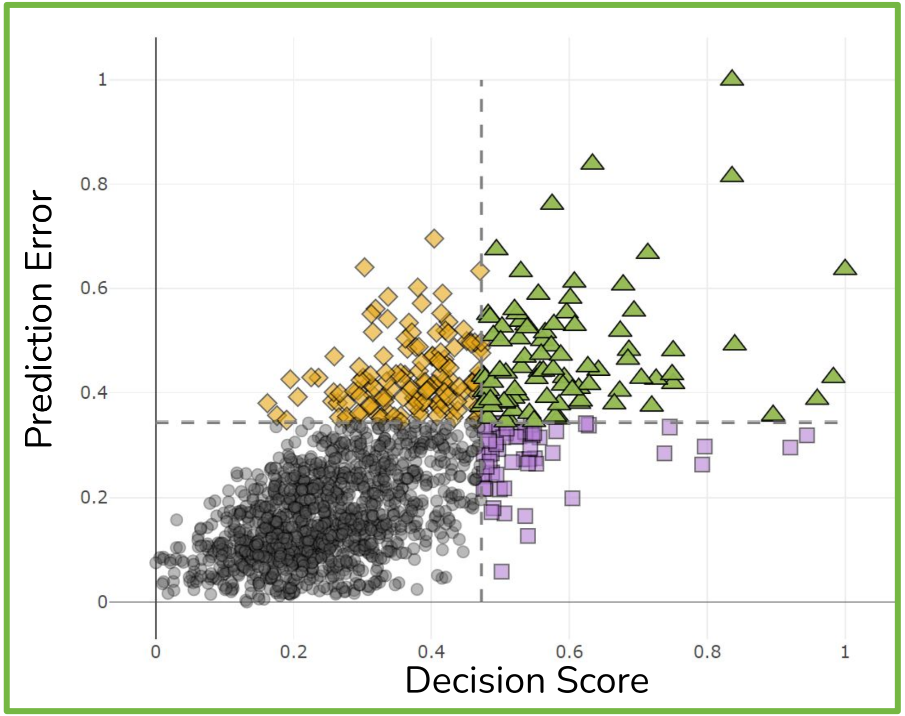

Counts	Metrics
TP: 91	Precision: 0.61
TN: 1216	Recall: 0.40
FN: 134	
FP: 59	

Results: Outlier Validation App

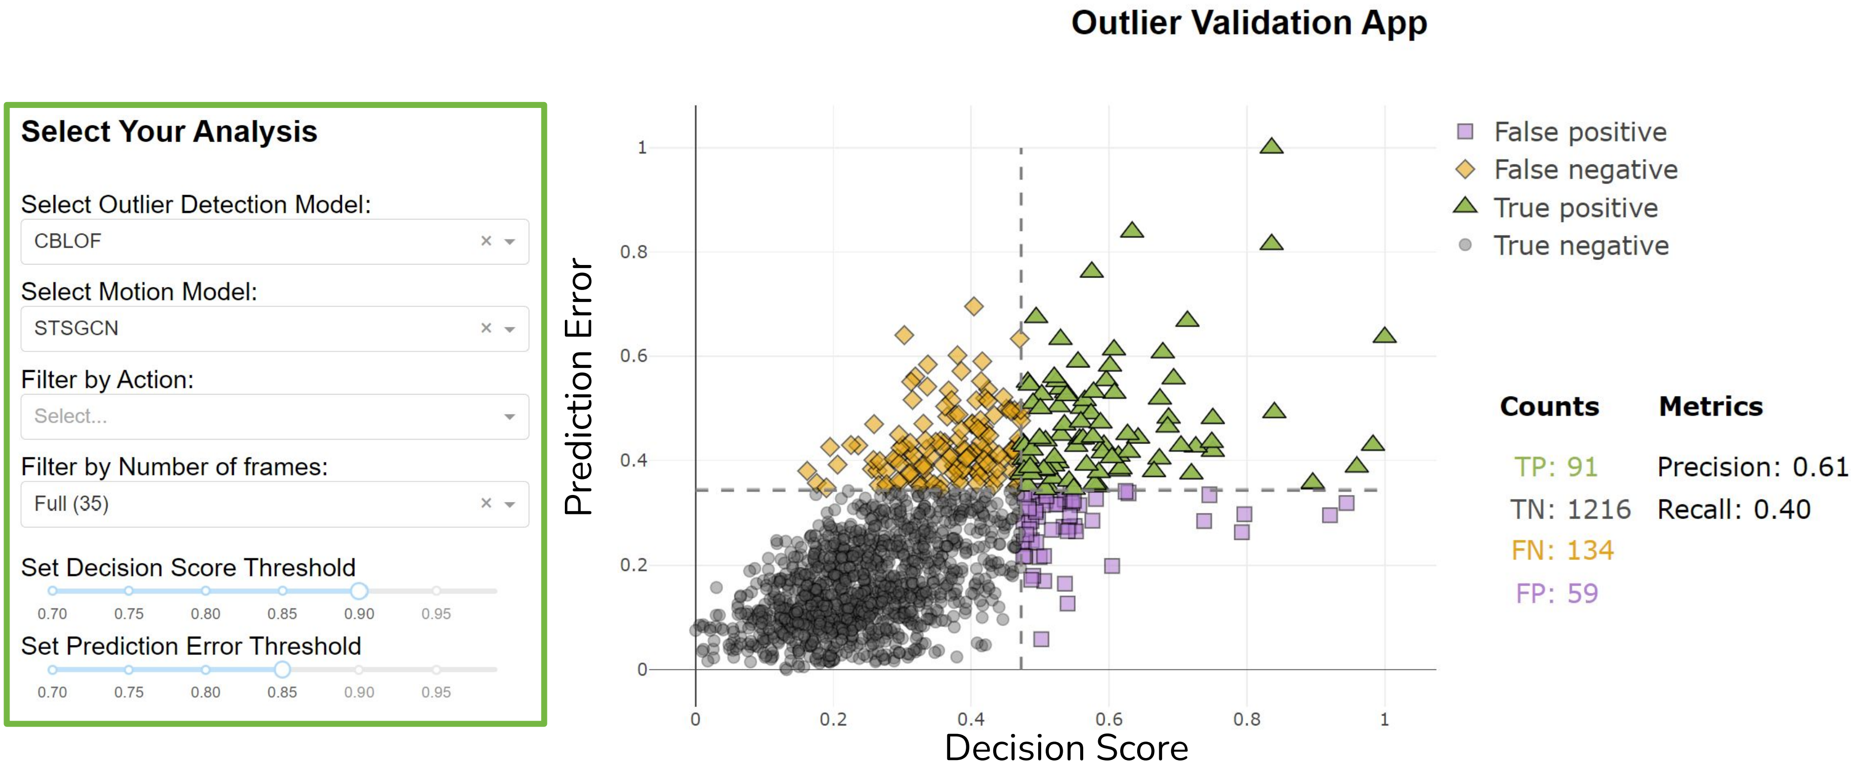

Results: Outlier Validation App

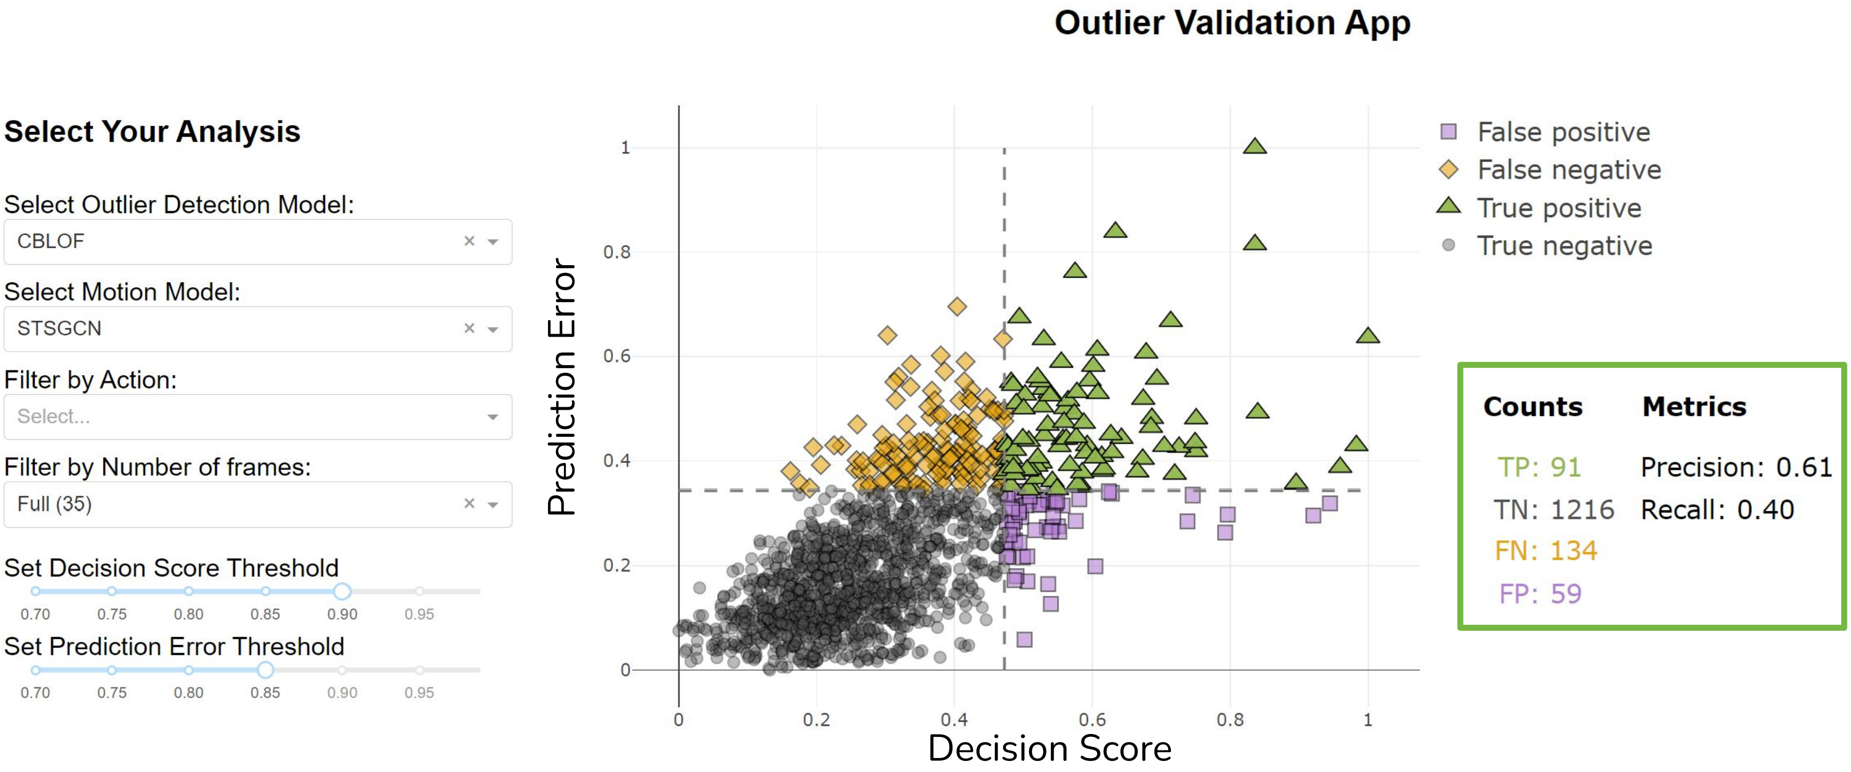

Results: Kinematic Comparison Toolkit

Inlier “walking”

Velocity

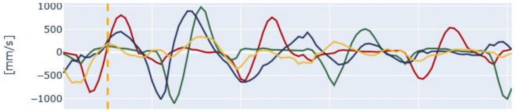

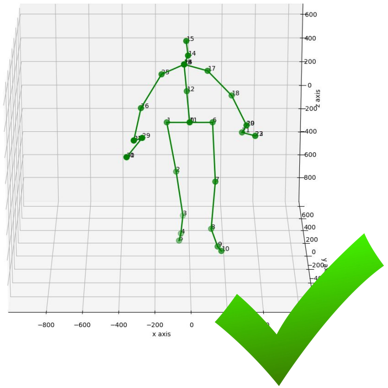

Outliers “walking”

Velocity

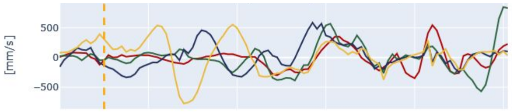

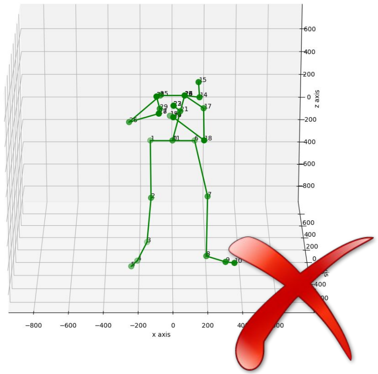

Velocity

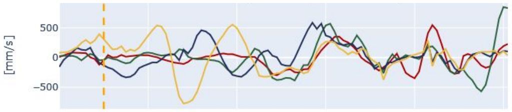

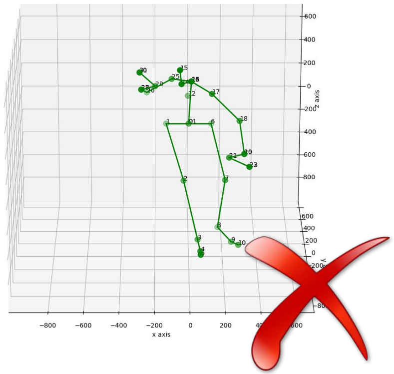

Summary & Outlook

Summary

→ **Business Case:**
Autonomous Driving

→ **Approach:**
Outlier detection models

→ **Results:**
Apps for outlier analysis

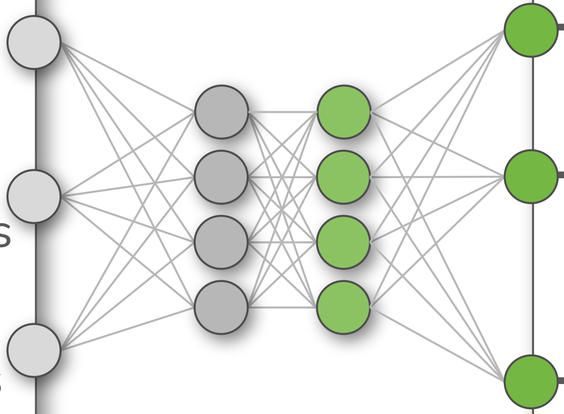

Outlook

→ **Expand:**
Generalize to other data

→ **Improve:**
Motion predictions models

→ **Utilize:**
Other project branches

One Step Ahead - Stay in touch

Don't be shy, just talk to us ;)

Jonas Voßemer

Alaa Elshorbagy

Vincent v. Zitzewitz

Our LinkedIn + App Demo here

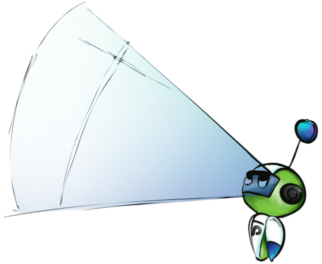

Backup

Agenda

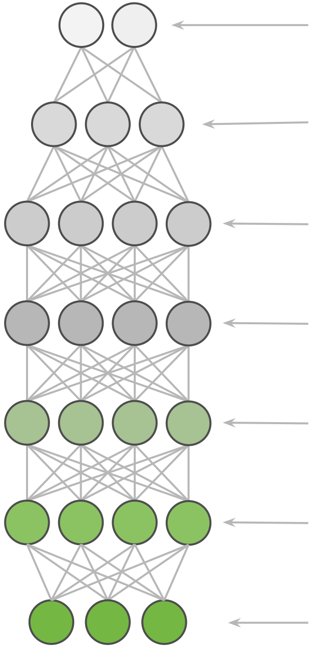

- Human Motion Prediction Applications
- QualityMinds Human Motion Prediction Project
- Objective: Data analysis and learning representation
- The Data we used
- Our Approach
- Key Results
- Outlook / Future Steps

Human Motion Prediction Applications

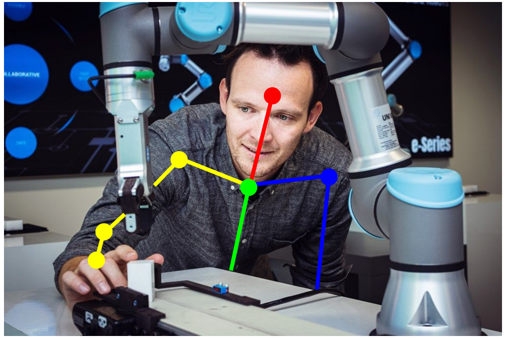

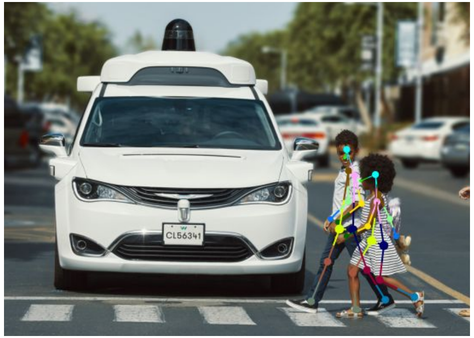

Sports Performance Analysis: Enhance athlete training programs | **Robotics & Manufacturing:** Optimize human-robot collaboration | **Autonomous Vehicles:** Improve pedestrian safety measures | **Healthcare Rehabilitation:** Guide physical therapy exercises | **Fitness Tracking:** Personalized exercise recommendations | **Security & Surveillance:** Anomaly detection in public spaces | **Ergonomics Design:** Workplace injury prevention planning | **Human-Computer Interaction:** Gesture-based interfaces advancement.| **Biomechanical Research:** Study human movement patterns

Input

Output (Prediction)

QM Human Motion Prediction Project

Motion prediction model

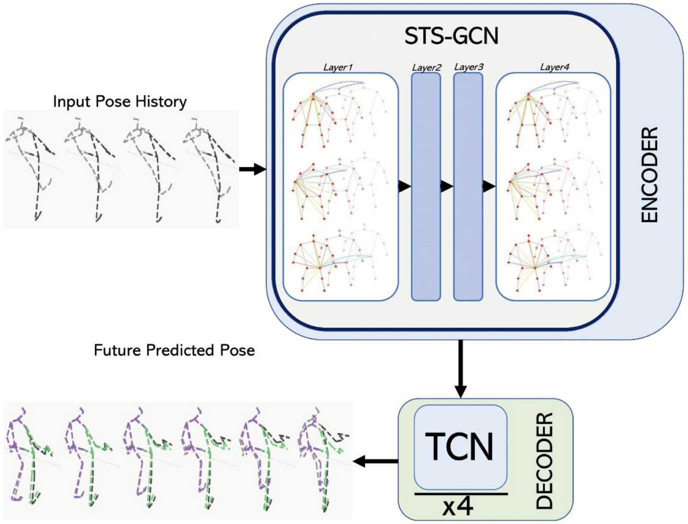

Data:
- Input: Sequence of 10 frames = 0.4 sec
- 10 frames x 32 joints x 3 coordinates

Model: STS-GCN
- **S**pace **T**ime **S**eparable -
Graph **C**onvolutional **N**etwork

Prediction:
- Output: Sequence of 25 frames = 1.0 sec

QM Human Motion Prediction Project: Branches

Five project branches

1. Data analysis (Learning representation)
2. Adversarial attacks (Robust learning)
3. Model interpretability
4. Efficient learning (Self-supervised Learning)
5. Simulation

Data analysis and learning representation

What is our **main task / subproject** in this project branch?

- **Find outlying motion sequences**
  - **Quantify the degree of anomaly of a motion sequence as a vector,**
    e.g., idx=35, action=”walking” → decision scores: [0.83, 0.92, 0.67, 0.89]
- Characterize outlying vs. inlying motion sequences
  - Describe relevant sequence characteristics, e.g., velocity or acceleration of specific joints visually or statistically

Data analysis and learning representation

Why is this important in the context of the overall project?

- **Motivating assumption** for our task
    - **Outlying motion sequence → high error in motion prediction models**
    - Improve motion prediction model
- Examples of an **outlying motion sequence**
    - Combining different actions: “eating” while “walking”
    - Sudden movements

Data: Different Datasets

Pose Prediction Datasets:

- **AMASS:** Motion capturing & in the wild recordings. Annotated with skeleton information
- **Human3.6M:** 3.6 million poses, motion capturing, 11 actors, 17 scenarios (walking, phoning, eating,...), 3D laser scans of the actors
- **3DPW:** Generated from real-world videos. Includes varying clothing, background, activities
- **CMU:** Motion capturing, 2D and 3D joint trajectories
- **ExPI:** Understand human emotions in human-robot interaction scenarios of human pushing the robot

Human3.6M: Labeling (although not always reliable), good as a generalizing approach, high variability

Data: H3.6M Dataset Structure

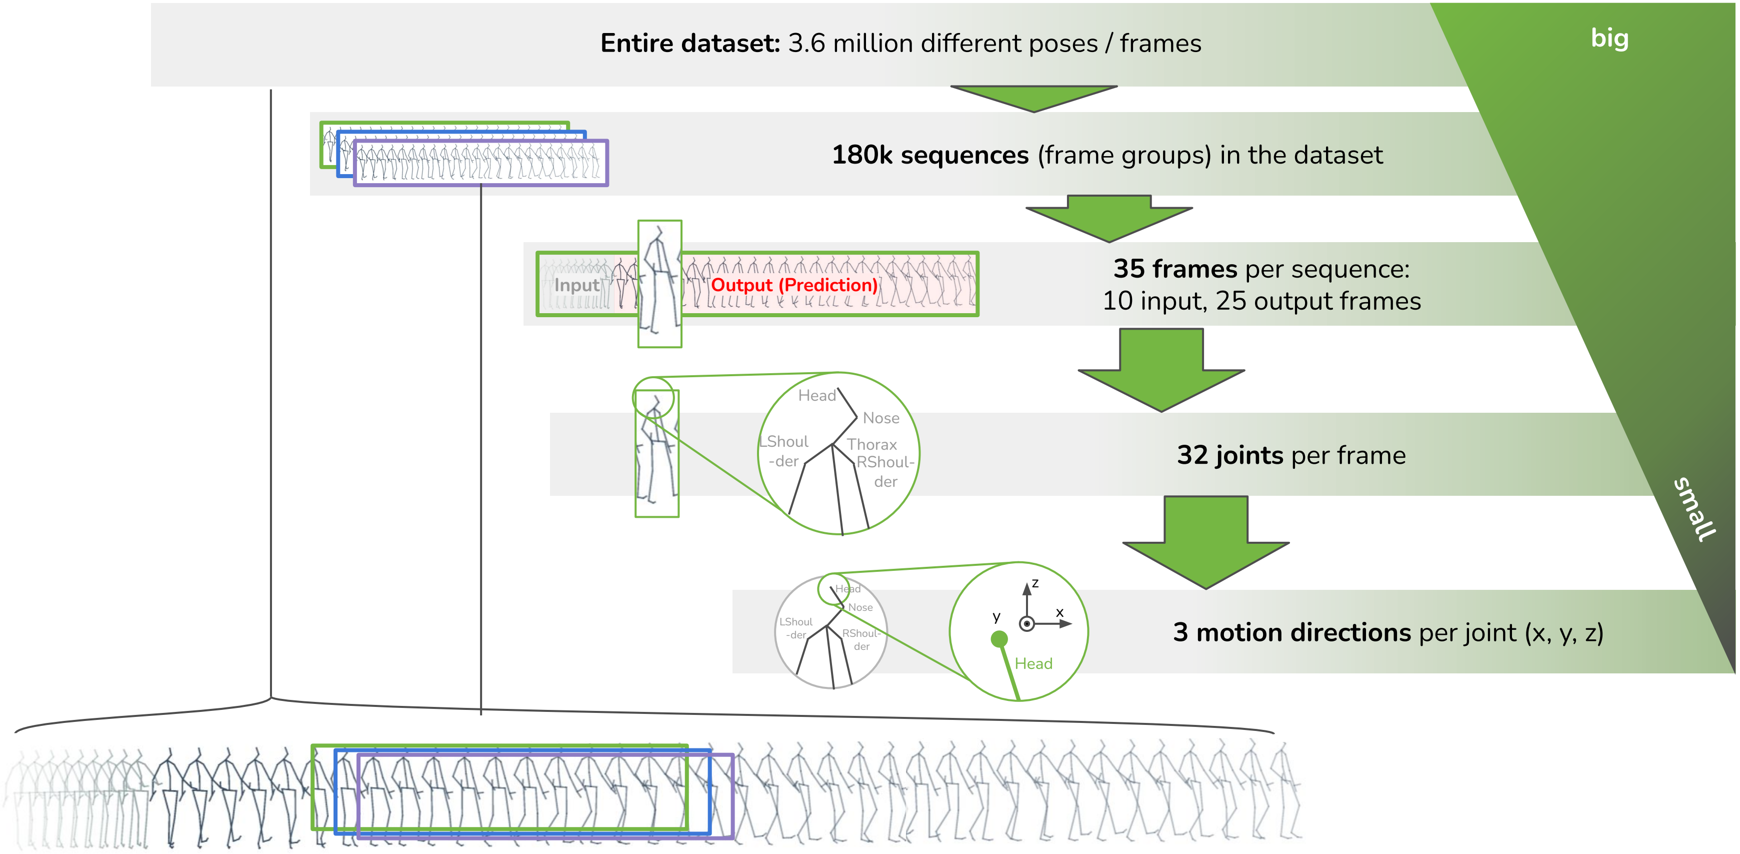

Data: Processing and Transformations

Raw Data

4D Array:
180k sequences
x 35 frames
x 32 joints
x 3 motion directions

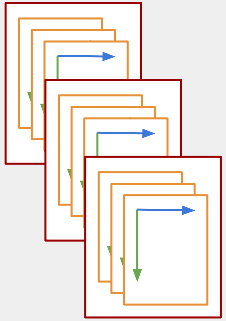

Data Processing

2D Array / Dataframe:
180k sequences
x (35 x 32 x 3)

3360 columns

180k rows

	idx	0	1	2	3	4	5	6
0	0	0.0	0.0	0.0	-132.948593	0.0	0.0	-71.220520
1	1	0.0	0.0	0.0	-132.948593	0.0	0.0	-70.484131
2	2	0.0	0.0	0.0	-132.948593	0.0	0.0	-69.480377
3	3	0.0	0.0	0.0	-132.948593	0.0	0.0	-71.031143
4	4	0.0	0.0	0.0	-132.948593	0.0	0.0	-71.140823

Data Transformation 1

Cartesian Coordinates

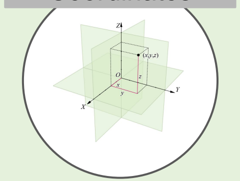

Spherical Coordinates

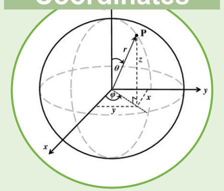

Data Transformation 2

S. Coordinates

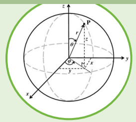

dt

Velocity in SC

dt

Acceleration in SC

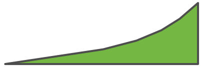

Methodology - Overview of our pipeline

How do we find outlying motion sequences?

Data representation	Select # of frames	Reduce # of dimensions (optional)	Scale data (optional)	Outlier detection method(s)	Decision score(s)
Cartesian coordinates - joint pos. (35 x 32 x 3) **Spherical coordinates** - joint pos. - velocity - acceleration	**Full:** All 35 frames **Input only:** First 10 frames **Output only:** Last 25 frames	**PCA** - 50 principal components - contain about 96% of variation in original space (35x32x3 = 3360 dim.)	**Standard scaler (z-score)** - Mean = 0, SD=1	**4 different unsupervised methods** - CBLOF - HDBSCAN - IFOREST - LOF	**Quantification of “anomaly”:** Higher values → more outlying Lower values → less outlying

Methodology - Outlier Detection

Main interest of QM: **Vector of normalized decision scores**

idx	action	CB-LOF	HDBSCAN	IFOREST	LOF
10	walking	0,85	0,76	0,59	0,80
84	eating	0,23	0,35	0,28	0,50
...	...	...	...	...	...

But why multiple outlier detection methods? → Each defines anomaly differently

→ **More diverse representation!**

Method	Which sequences are considered anomalous, i.e. have a higher decision score?
CB-LOF	Separating into small/large clusters (k-means) → Sequences with a high distance to the nearest large cluster
HDBSCAN	Sequences not belonging to any density-based cluster or being at the edge of a density-based cluster
IFOREST	Sequences that are easily isolated in a random tree, i.e. a low number of splits required (averaged over a forest)
LOF	Sequences with significantly lower local density compared to their k nearest neighbors

Methodology - Overview of our pipeline

How do we find outlying motion sequences?

Data representation	Select # of frames	Reduce # of dimensions (optional)	Scale data (optional)	Outlier detection method(s)
Cartesian coordinates - joint pos. (35 x 32 x 3) **Spherical coordinates** - joint pos. - velocity - acceleration	**Full:** All 35 frames **Input only:** First 10 frames **Output only:** Last 25 frames	**PCA** - 50 principal components - contain about 96% of variation in original space (35x32x3 = 3360 dim.)	**Standard scaler (z-score)** - Mean = 0, SD=1	**4 different unsupervised methods** - CBLOF - HDBSCAN - IFOREST - LOF

Decision score(s)

Prediction errors from motion models

Validation of outliers

Methodology - Validation

Problem: Unsupervised approach lacks ground truth - Does it work?

- **Explorative:** Compare outlying / inlying sequences on relevant characteristics
  - Visualize sequences, e.g. motion GIFs, velocity/acceleration of joints
  - Characterize them statistically, e.g., average velocity/acceleration of joints
- **Artificial Metrics:** Compare outlying / inlying sequences using prediction errors
  - Merge **prediction errors** from **two motion models** (STS-GCN, motionmixer) to Data

Methodology - Validation

Prediction error distribution from motion models →
"True outliers"

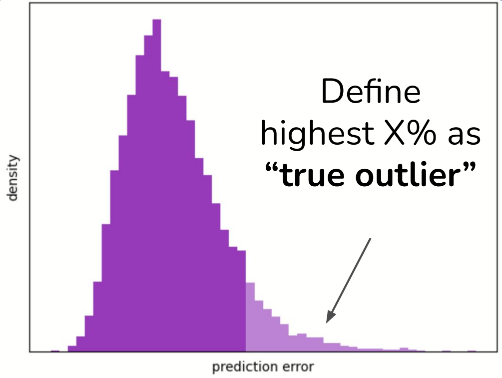

Decision score distribution outlier detection methods →
"Predicted outliers"

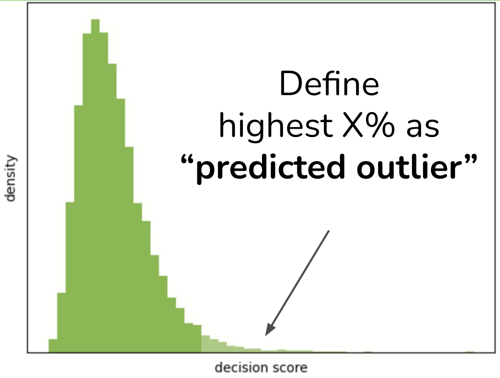

Calculate **artificial precision and recall metrics**
based on "true outlier" and "predicted outlier" definitions

Methodology - Validation

- **Artificial “precision” and “recall”**
    - Precision = TP / TP+FP
        - Rate of predicted outlier sequences with high prediction error
    - Recall = TP / TP+FN
        - Rate of high error sequences captured by predicted outliers
- **Flexibility in definitions due to:**
    - Variable thresholds for highest X%
    - Use single outlier detection method or combinations of these
    - Use single motion model or combination of two motion models

Results - Out of Distribution (OOD) sequences

OOD Methods

- CBLOF
- IForest
- HDBScan
- LOF

Outlier detection:
Threshold: 0.9

MPJPE error: minimum error of STS-GCN and motionmixer

Detected outliers projected in 2D plane (using t-SNE)

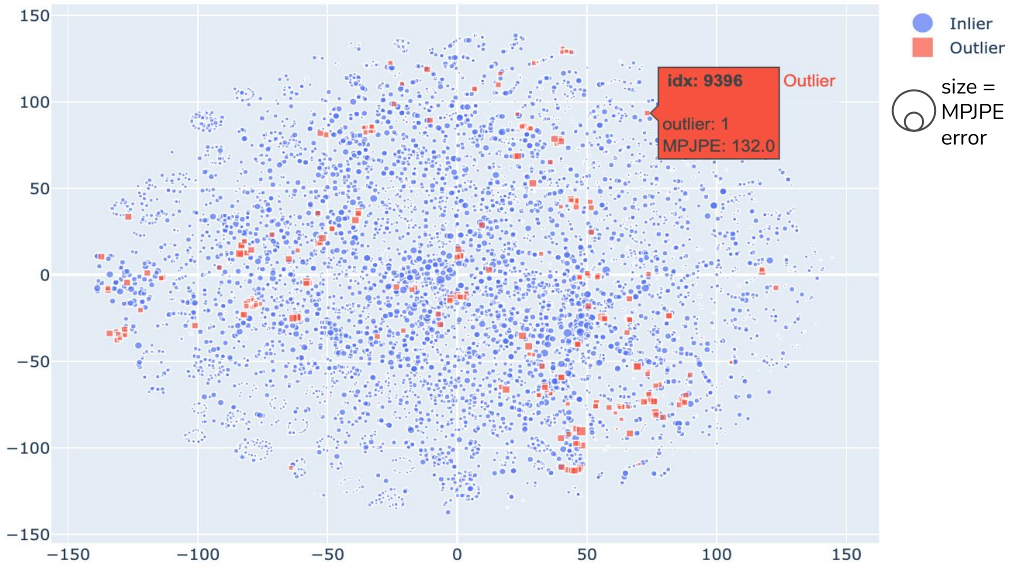

Results - Out of Distribution Sample Detection

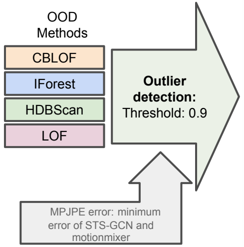

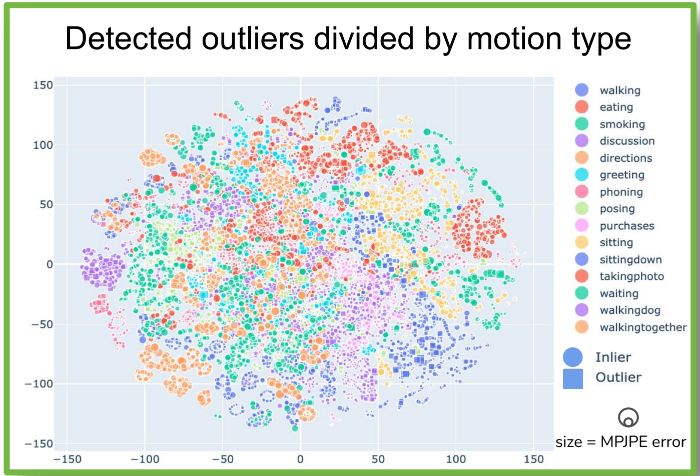

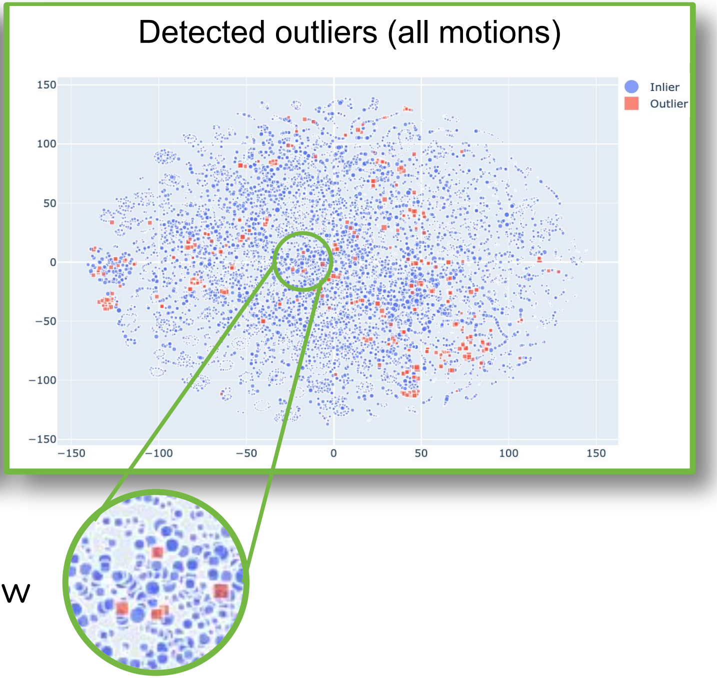

Detail view

Results - Analysis Toolkit

Comparison of outliers based on motion model and decision scores

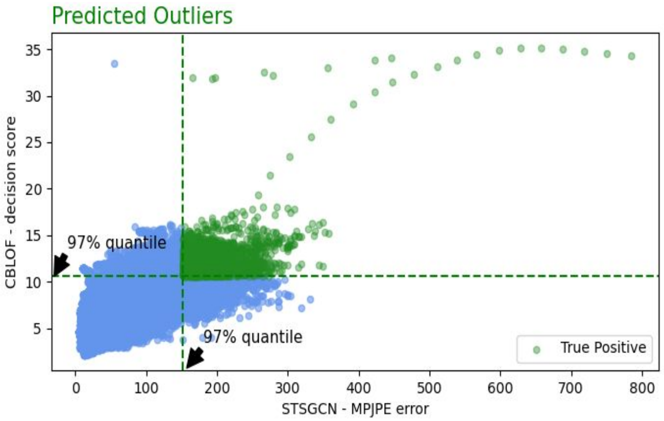

Top 3 % of MPJPE Errors in STSGCN model and decision scores of CBLOF (anomaly detection model)

Comparison of motion prediction models vs OOD detection methods

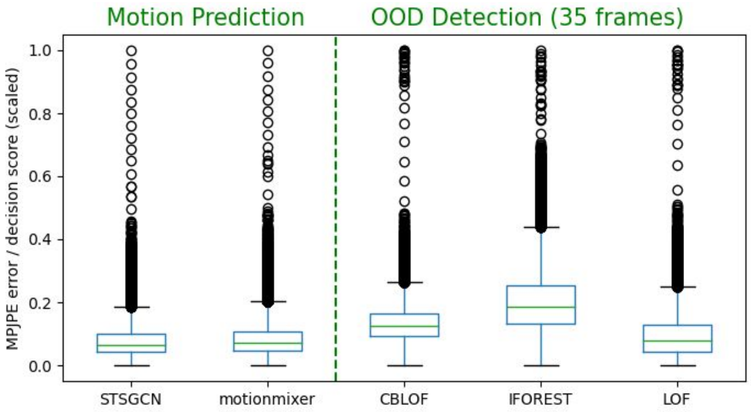

Results - Analysis Toolkit

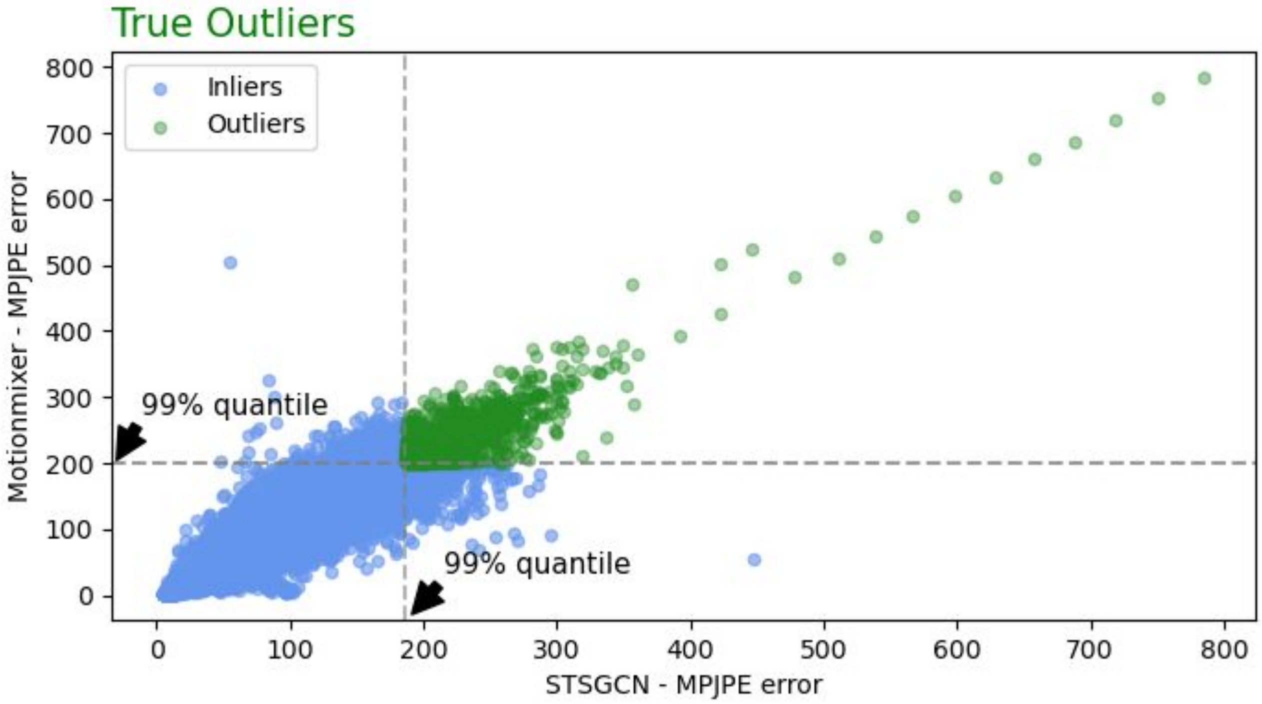

Results - Analysis Toolkit Recall & Precision

	motion_model_th:0.97	OOD_model_th:0.9	mode	recall	precision	recall_&	precision_&	recall_or	precision_or
0	STSGCN	CBLOF	full	0,71	0,21	0,50	0,30	0,80	0,15
1	STSGCN	IFOREST	full	0,61	0,18				
2	STSGCN	HDBSCAN	full	0,42	0,13				
3	STSGCN	LOF	full	0,65	0,19				
4	STSGCN	CBLOF	Ti_10	0,32	0,10	0,21	0,14	0,44	0,08
5	STSGCN	IFOREST	Ti_10	0,31	0,09				
6	STSGCN	HDBSCAN	Ti_10	0,26	0,08				
7	STSGCN	LOF	Ti_10	0,34	0,10				
8	STSGCN	CBLOF	To_25	0,74	0,22	0,52	0,31	0,84	0,16
9	STSGCN	IFOREST	To_25	0,65	0,19				
10	STSGCN	HDBSCAN	To_25	0,60	0,18				
11	STSGCN	LOF	To_25	0,68	0,21				
12	motionmixer	CBLOF	full	0,71	0,21	0,51	0,31	0,80	0,15
13	motionmixer	IFOREST	full	0,62	0,19				
14	motionmixer	HDBSCAN	full	0,43	0,13				
15	motionmixer	LOF	full	0,66	0,20				
16	motionmixer	CBLOF	Ti_10	0,34	0,10	0,22	0,15	0,47	0,09
17	motionmixer	IFOREST	Ti_10	0,33	0,10				
18	motionmixer	HDBSCAN	Ti_10	0,27	0,08				
19	motionmixer	LOF	Ti_10	0,36	0,11				
20	motionmixer	CBLOF	To_25	0,74	0,22	0,52	0,32	0,84	0,16
21	motionmixer	IFOREST	To_25	0,64	0,19				
22	motionmixer	HDBSCAN	To_25	0,61	0,18				
23	motionmixer	LOF	To_25	0,69	0,21				

Results - Analysis Toolkit Recall & Precision

```
general_analysis(metrics_df,model_threshold = 0.90, error_threshold= 0.97)
```

	modes_mt=0.9_et=0.97	recall_&	precision_&	recall_or	precision_or
0	full	0,44	0,34	0,83	0,08
1	Ti_10	0,20	0,13	0,50	0,05
2	To_25	0,54	0,29	0,89	0,10

```
general_analysis(metrics_df,model_threshold = 0.95, error_threshold= 0.90)
```

	modes_mt=0.95_et=0.9	recall_&	precision_&	recall_or	precision_or
0	full	0,16	0,74	0,48	0,37
1	Ti_10	0,10	0,44	0,28	0,23
2	To_25	0,20	0,70	0,50	0,48

Results - Analysis Toolkit Sph. Velocity & Acceleration

Velocity in Spherical Coordinates

Acceleration in Spherical Coordinates

Sequence Analysis (e.g. walking sequence 100)

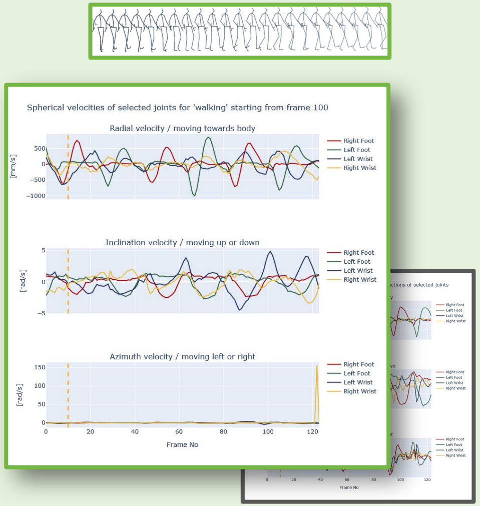

Motion Analysis (e.g. greeting)

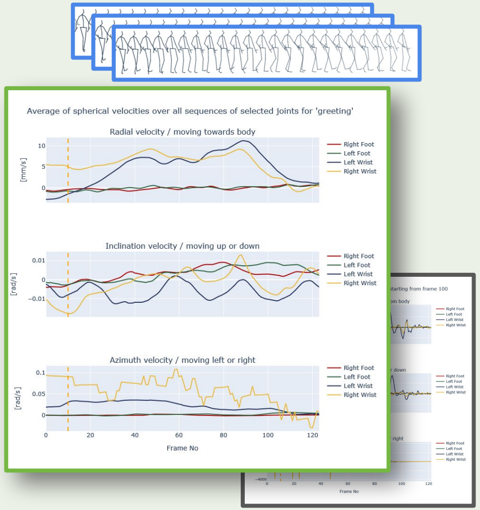

All data analysis

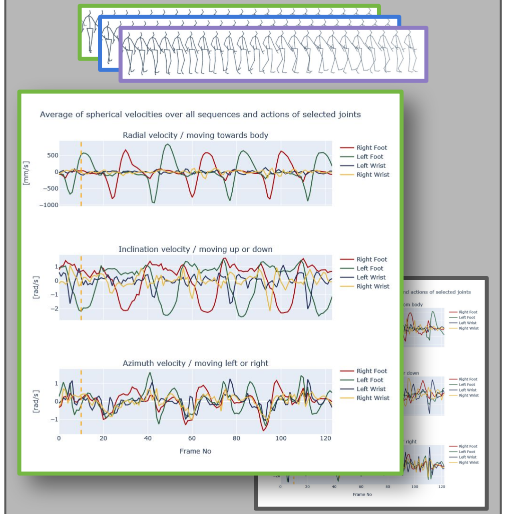

Results - Analysis Toolkit Sph. Velocity & Acceleration

Velocity in Spherical Coordinates

Acceleration in Spherical Coordinates

Inlier Analysis (e.g. walking sequence 3116)

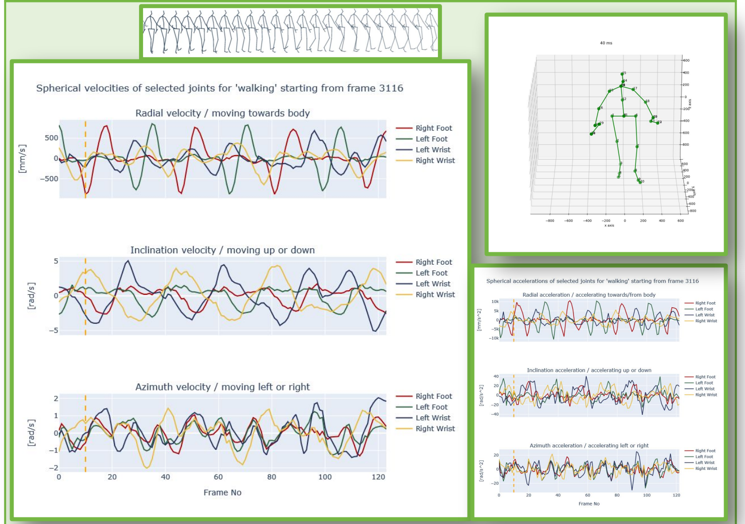

Outlier Analysis (e.g. smoking frame 9396)

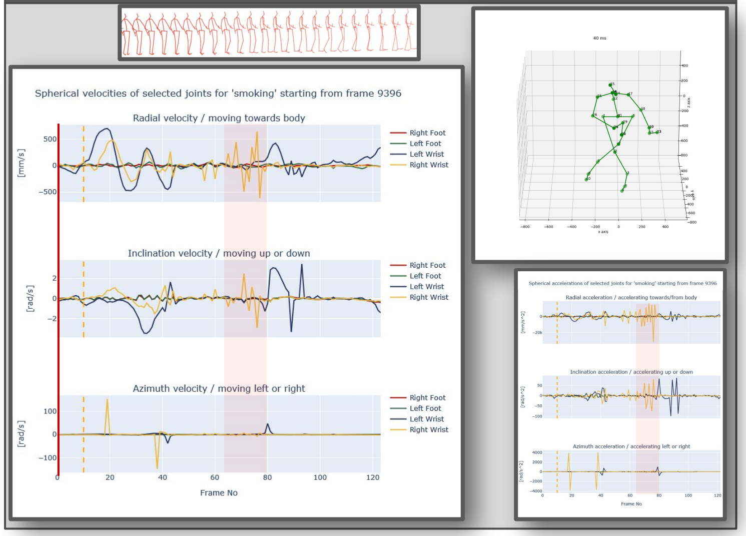

Outlook

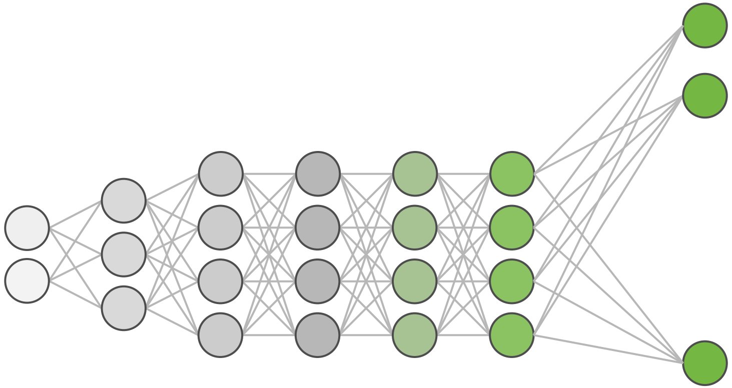

Expand: Generalize to other public datasets

Improve: Use the vector representations to enhance the performance of state of the art motion predictions models

- Regularize the loss function (test the OOD methods validity)
- Do a parallel / meta network to modify the weights of the original motion prediction models

Utilize: Use the representation in the other four stages of the project

- Perform adversarial attacks based on the found outliers
- Enhance model Interpretability
- Support the self learning of the model
- Adapt the simulation branch

Results - OOD (methods combined)

T-SNE plot

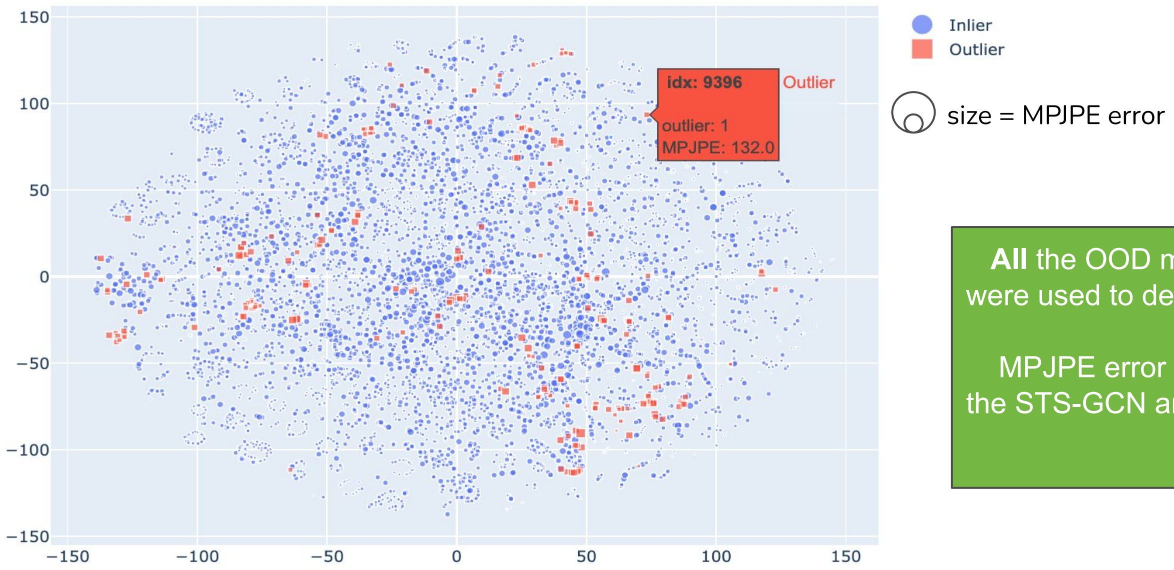

All the OOD methods (threshold 0.90) were used to detect outliers.

MPJPE error is the minimum error of the STS-GCN and motionmixer.

Results - Out of Distribution Sample Detection

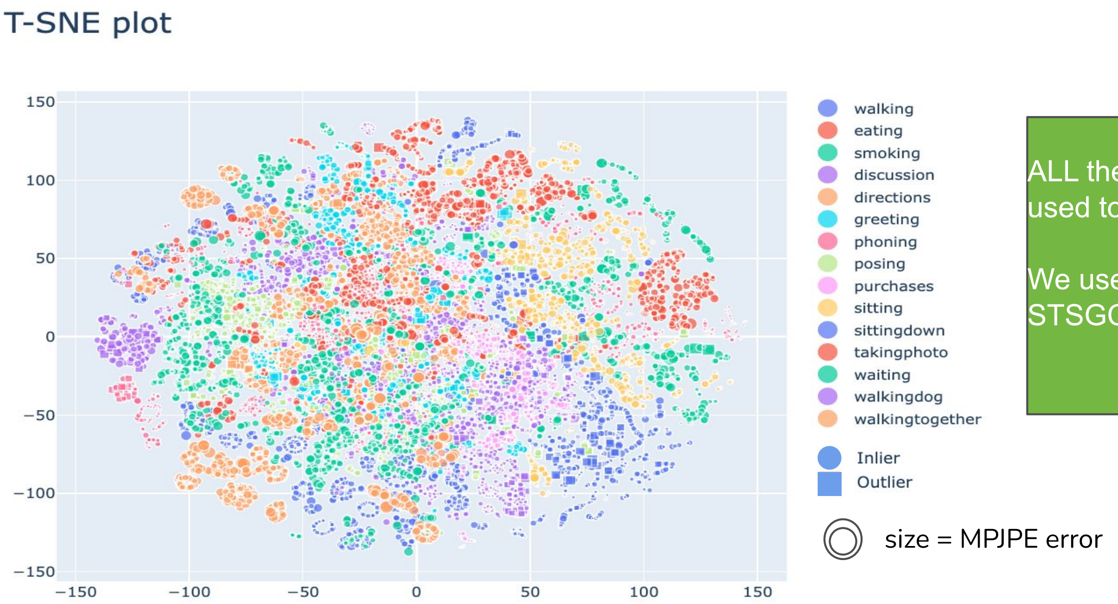

ALL the OOD methods (threshold 0.90) is used to detect outliers.

We use the minimum MPJPE error of the STSGCN and motionmixer. .

Inliers and Outliers “walking” sequence

Inlier

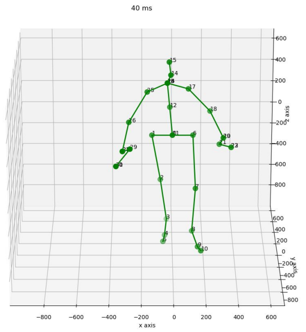

Outliers

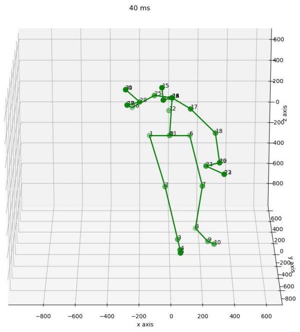

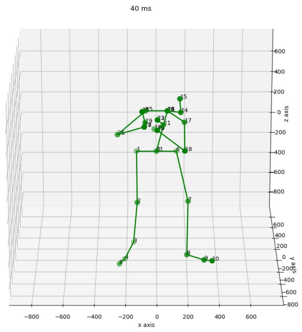

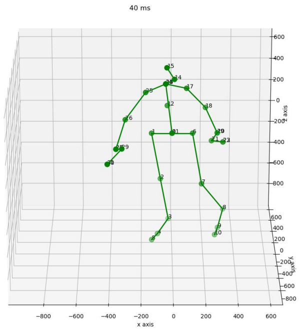

Teamwork makes the dream work

Teamwork makes
the dream work

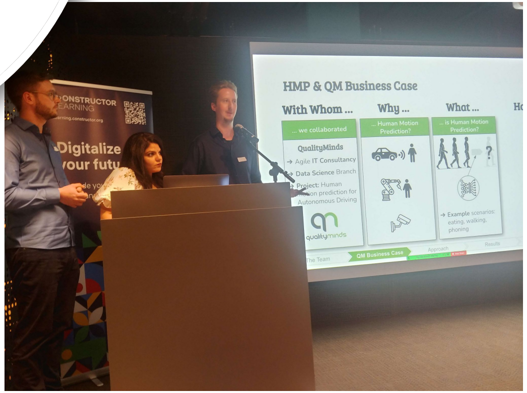

Frame set 100 (frame 100 + frame length, e.g. +75 frames)

Frame set 2 (frame 205 + frame length, e.g. +75 frames)

Metric numbers:

*

Viz of the pose (take from demo file)

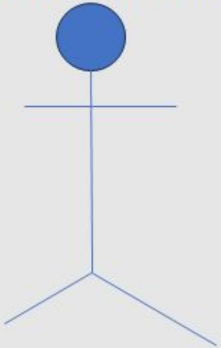

Scatter plot:

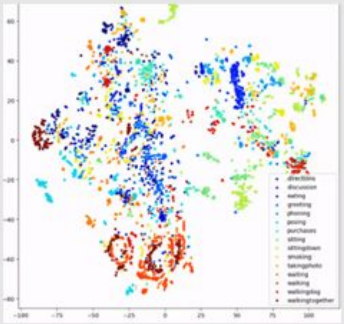

Velocity all outliers

Velocity all inliers

Velocity one sample

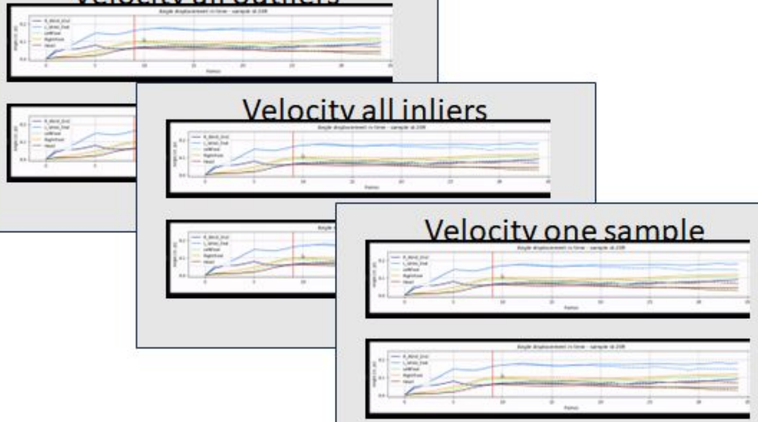

Decision Score: Probability of the point being an outlier / Error of point

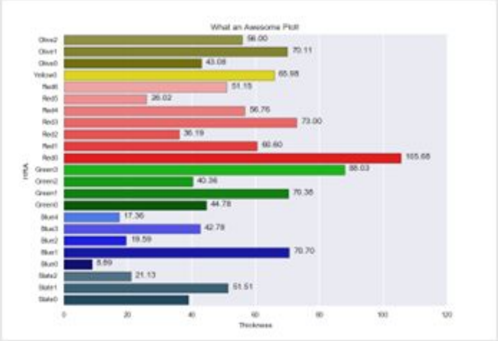

Frame set 100 (frame 100 + frame length, e.g. +75 frames)

Frame set 2 (frame 205 + frame length, e.g. +75 frames)

Metric numbers:

*

Viz of the pose (take from demo file)

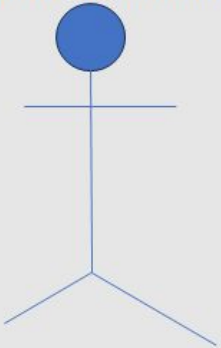

Scatter plot:

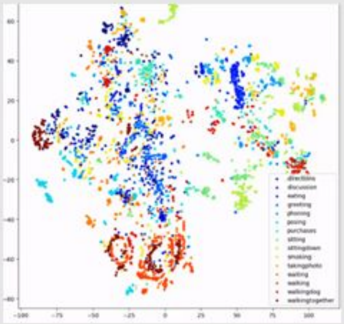

Velocity all outliers

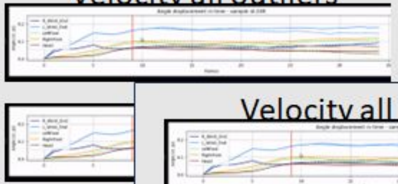

Velocity all inliers

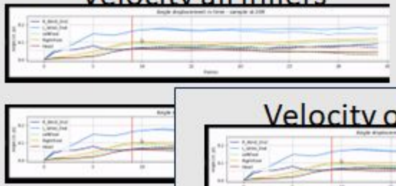

Velocity one sample

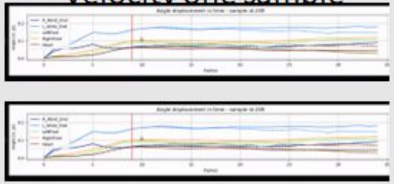

Decision Score: Probability of the point being an outlier / Error of point

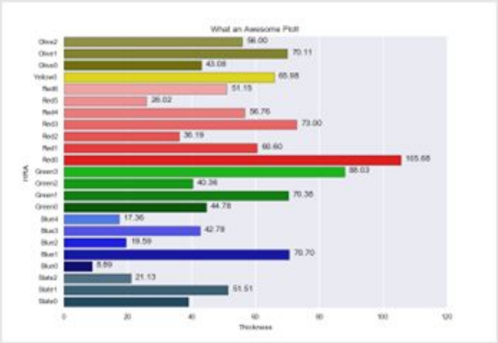

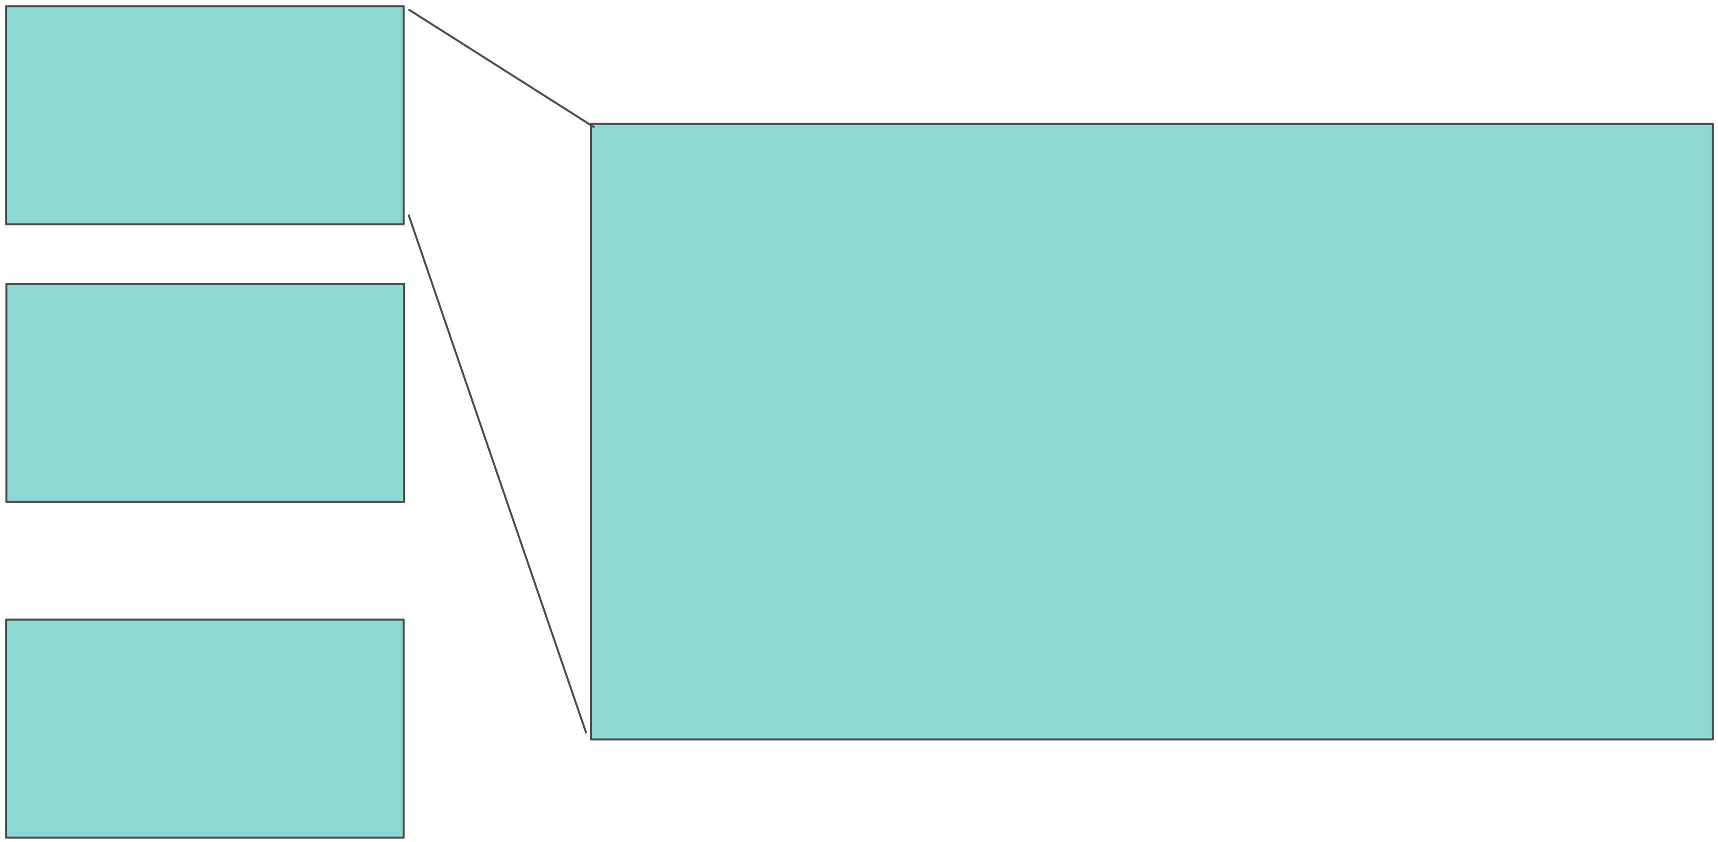

Intro us

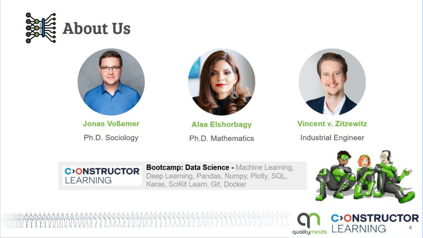

Intro QM + Business Case

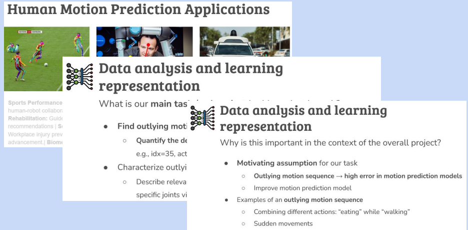

Our Approach

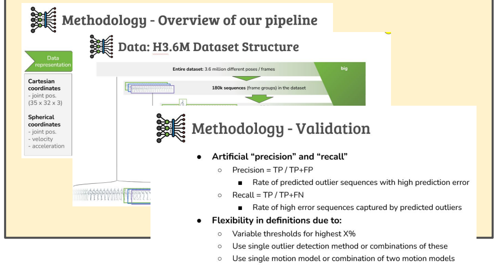

Results: OOD Samples

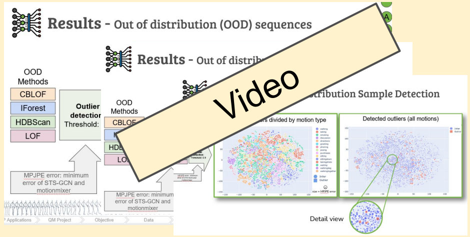

Results: Analysis toolkit

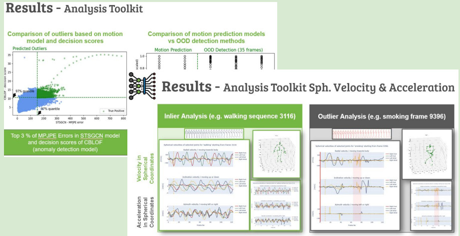

Outlook

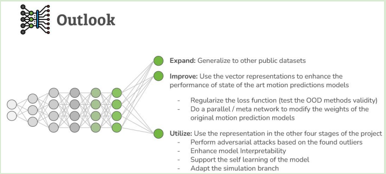

Intro us again (w linkedin)

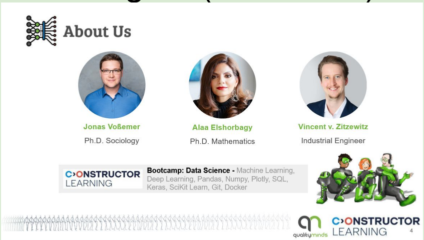

Presentation structure:

1. Introduce yourselves (w/ backgrounds)
2. Introduce the client/company (What do they do? What are their goals; what do they want out of the project?)
3. Your approach
4. Results & conclusion (with demo if applicable)
5. Future considerations
6. Reiterate again who you were with contact information (e.g. a link to your LinkedIn)

Our Approach

Data Structure

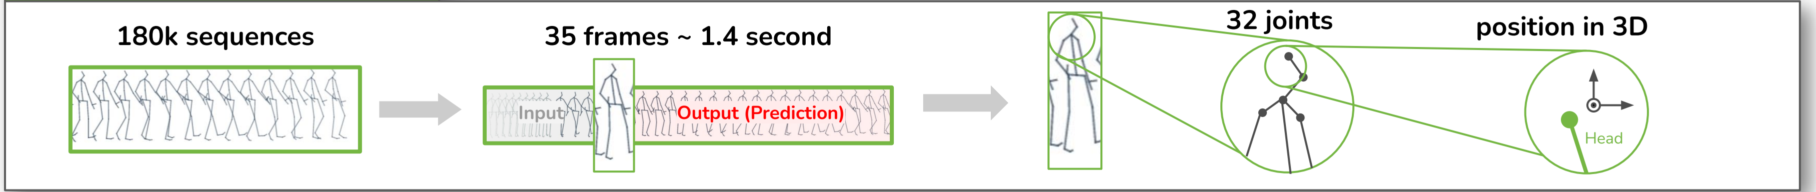

Methodology

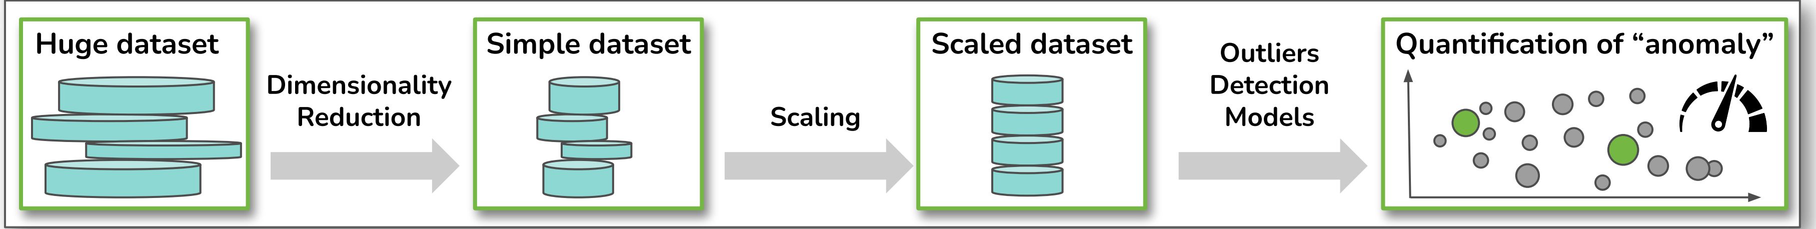

Validation